

```
In [23]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib as mpl
import matplotlib.pyplot as plt
# For randomized data splitting
from sklearn.model_selection import train_test_split

# To build linear regression model
import statsmodels.api as sm
# To check model performance
from sklearn.metrics import mean_absolute_error, mean_squared_error
import pickle
```

```
In [24]: df = pd.read_csv('/Users/haoyuechang/Desktop/auto-mpg.csv')
#cData = pd.read_csv("auto-mpg.csv")
```

```
: df = pd.read_csv('/Users/haoyuechang/Desktop/auto-mpg.csv')
#cData = pd.read_csv("auto-mpg.csv")
```

```
: #Data processing
df.info()
#398 entries, 9 columns
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 398 entries, 0 to 397
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   mpg              398 non-null   float64
1   cylinders        398 non-null   int64
2   displacement     398 non-null   float64
3   horsepower       398 non-null   object
4   weight           398 non-null   int64
5   acceleration     398 non-null   float64
6   model year       398 non-null   int64
7   origin           398 non-null   int64
8   car name         398 non-null   object
dtypes: float64(3), int64(4), object(2)
memory usage: 28.1+ KB
```

```
: #drop car name since it is not useful to analyze
df1 = df.drop(["car name"], axis=1)
```

```

: hpIsDigit = pd.DataFrame(
    df1.horsepower.str.isdigit()
) # if the string is made of digits store True else False

# print the entries where isdigit = False
df1[hpIsDigit["horsepower"] == False]

```

```

:
      mpg  cylinders  displacement  horsepower  weight  acceleration  model year  origin
32  25.0         4         98.0           ?    2046         19.0         71         1
126  21.0         6        200.0           ?    2875         17.0         74         1
330  40.9         4         85.0           ?    1835         17.3         80         2
336  23.6         4        140.0           ?    2905         14.3         80         1
354  34.5         4        100.0           ?    2320         15.8         81         2
374  23.0         4        151.0           ?    3035         20.5         82         1

```

```

: df1= df1.replace("?", np.nan)
df1[hpIsDigit["horsepower"] == False]

```

```

:
      mpg  cylinders  displacement  horsepower  weight  acceleration  model year  origin
32  25.0         4         98.0         NaN    2046         19.0         71         1
126  21.0         6        200.0         NaN    2875         17.0         74         1
330  40.9         4         85.0         NaN    1835         17.3         80         2
336  23.6         4        140.0         NaN    2905         14.3         80         1
354  34.5         4        100.0         NaN    2320         15.8         81         2
374  23.0         4        151.0         NaN    3035         20.5         82         1

```

```
] : df1.median()
```

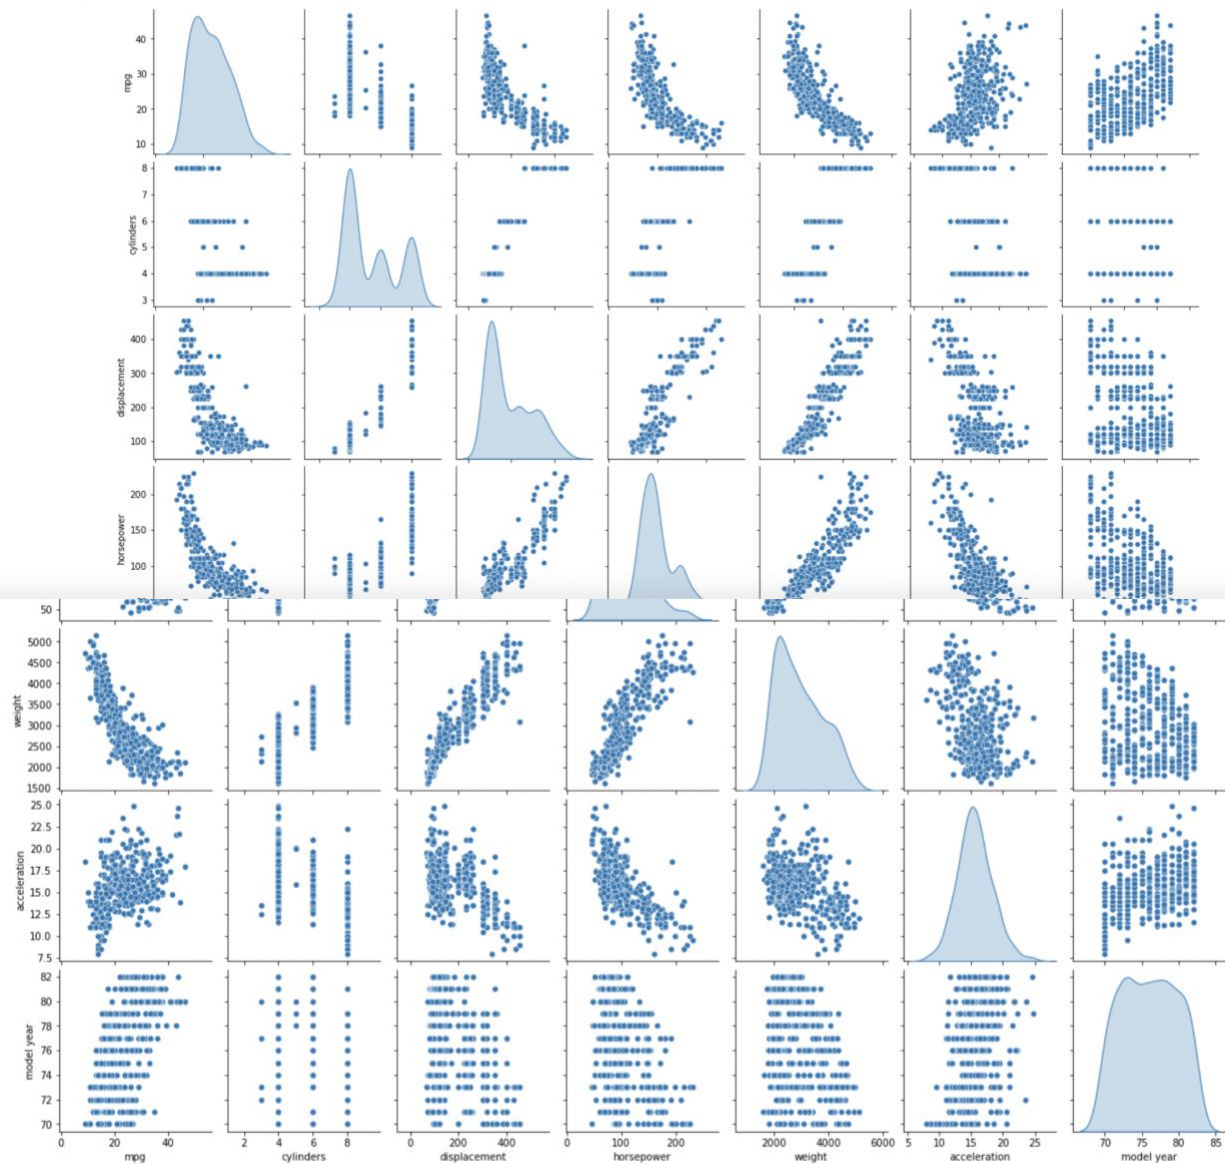
```
] : mpg          23.0  
    cylinders     4.0  
    displacement 148.5  
    horsepower    93.5  
    weight       2803.5  
    acceleration  15.5  
    model year    76.0  
    origin        1.0  
    dtype: float64
```

```
] : # Let's replace the missing values with median values of the columns.  
# Note that we do not need to specify the column names below.  
# Every column's missing value is replaced with that column's median respectively  
  
medianFiller = lambda x: x.fillna(x.median())  
df1 = df1.apply(medianFiller, axis=0)
```

```
] : # let's convert the horsepower column from object type to float type  
df1["horsepower"] = df1["horsepower"].astype(float)
```

```
In [32]: #Bivariate
df_attr = df1.iloc[:, 0:7]
sns.pairplot(
    df_attr, diag_kind="kde")
```

Out[32]: <seaborn.axisgrid.PairGrid at 0x7fb13015ca00>



```
In [33]: # drop_first=True will drop one of the three origin columns
#creat dummy variables
df2 = pd.get_dummies(df1, columns=["origin"], drop_first=True)
df2.head()
```

```
Out[33]:
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	model year	origin_2	origin_3
0	18.0	8	307.0	130.0	3504	12.0	70	0	0
1	15.0	8	350.0	165.0	3693	11.5	70	0	0
2	18.0	8	318.0	150.0	3436	11.0	70	0	0
3	16.0	8	304.0	150.0	3433	12.0	70	0	0
4	17.0	8	302.0	140.0	3449	10.5	70	0	0

```
In [34]: #SPLIT DATA
# independent variables
X = df2.drop(["mpg"], axis=1)
# dependent variable
y = df2[["mpg"]]
```

```
In [35]: X_train, X_test, y_train, y_test = train_test_split(
        X, y, test_size=0.30, random_state=1)
```

```
|: print(X_train.head())
```

	cylinders	displacement	horsepower	weight	acceleration	model year	\
350	4	105.0	63.0	2215	14.9	81	
59	4	97.0	54.0	2254	23.5	72	
120	4	121.0	112.0	2868	15.5	73	
12	8	400.0	150.0	3761	9.5	70	
349	4	91.0	68.0	1985	16.0	81	

	origin_2	origin_3
350	0	0
59	1	0
120	1	0
12	0	0
349	0	1

```
|: print(X_test.head())
```

	cylinders	displacement	horsepower	weight	acceleration	model year	\
174	6	171.0	97.0	2984	14.5	75	
359	4	141.0	80.0	3230	20.4	81	
250	8	318.0	140.0	3735	13.2	78	
274	5	131.0	103.0	2830	15.9	78	
283	6	232.0	90.0	3265	18.2	79	

	origin_2	origin_3
174	0	0
359	1	0
250	0	0
274	1	0
283	0	0

```
|: olsmod = sm.OLS(y_train, X_train)
   olsres = olsmod.fit()
```



```
: print(olsres.summary())
```

```

                        OLS Regression Results
=====
Dep. Variable:          mpg      R-squared (uncentered):          0.980
Model:                  OLS      Adj. R-squared (uncentered):        0.980
Method:                 Least Squares      F-statistic:          1689.
Date:                   Mon, 26 Jun 2023      Prob (F-statistic):      1.04e-225
Time:                   16:45:30      Log-Likelihood:         -741.29
No. Observations:       278      AIC:          1499.
Df Residuals:           270      BIC:          1528.
Df Model:                8
Covariance Type:        nonrobust
=====
                        coef      std err          t      P>|t|      [0.025      0.975]
-----
cylinders              -0.6654      0.427      -1.558      0.120      -1.506      0.175
displacement            0.0284      0.010       2.747      0.006       0.008      0.049
horsepower             -0.0443      0.016      -2.845      0.005      -0.075     -0.014
weight                -0.0067      0.001      -7.945      0.000      -0.008     -0.005
acceleration          -0.1202      0.110      -1.090      0.277      -0.337      0.097
model_year              0.6229      0.029     21.627      0.000       0.566      0.680
origin_2                2.3585      0.699       3.372      0.001       0.982      3.735
origin_3                2.1483      0.697       3.083      0.002       0.777      3.520
=====
Omnibus:                27.282      Durbin-Watson:          2.241
Prob(Omnibus):           0.000      Jarque-Bera (JB):        42.908
Skew:                    0.609      Prob(JB):                4.82e-10
Kurtosis:                4.490      Cond. No.                1.24e+04
=====
```

Notes:

- [1] R^2 is computed without centering (uncentered) since the model does not contain a constant.
- [2] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [3] The condition number is large, 1.24e+04. This might indicate that there are strong multicollinearity or other numerical problems.

```
# Saving model to disk
pickle.dump(olsres, open('model.pkl', 'wb'))
```

```
# Loading model to compare the results
model = pickle.load(open('model.pkl', 'rb'))
print(model.predict([[8,307.0,130.0,3504,12.0,70,0,0]]))

[16.27077424]
```

```
!pip3 install flask==2.2.2
```

```
Collecting flask==2.2.2
  Using cached Flask-2.2.2-py3-none-any.whl (101 kB)
Requirement already satisfied: Werkzeug>=2.2.2 in /opt/anaconda3/lib/python3.9/site-packages (from flask==2.2.2) (2.3.6)
Requirement already satisfied: Jinja2>=3.0 in /opt/anaconda3/lib/python3.9/site-packages (from flask==2.2.2) (3.1.2)
Requirement already satisfied: importlib-metadata>=3.6.0 in /opt/anaconda3/lib/python3.9/site-packages (from flask==2.2.2) (4.8.1)
Requirement already satisfied: itsdangerous>=2.0 in /opt/anaconda3/lib/python3.9/site-packages (from flask==2.2.2) (2.1.2)
Requirement already satisfied: click>=8.0 in /opt/anaconda3/lib/python3.9/site-packages (from flask==2.2.2) (8.1.3)
Requirement already satisfied: zipp>=0.5 in /opt/anaconda3/lib/python3.9/site-packages (from importlib-metadata>=3.6.0->flask==2.2.2) (3.6.0)
Requirement already satisfied: MarkupSafe>=2.0 in /opt/anaconda3/lib/python3.9/site-packages (from Jinja2>=3.0->flask==2.2.2) (2.1.2)
Installing collected packages: flask
  Attempting uninstall: flask
    Found existing installation: Flask 2.3.2
    Uninstalling Flask-2.3.2:
      Successfully uninstalled Flask-2.3.2
Successfully installed flask-2.2.2
```

```
import numpy as np
from flask import Flask, request, render_template
import pickle

app = Flask(__name__)
model = pickle.load(open('model.pkl', 'rb'))

@app.route('/')
def home():
    return render_template('index.html')

@app.route('/predict', methods=['POST'])
def predict():
    """
    For rendering results on HTML GUI
    """
    int_features = [int(x) for x in request.form.values()]
    final_features = [np.array(int_features)]
    prediction = model.predict(final_features)

    output = round(prediction[0], 2)

    return render_template('index.html', prediction_text='a car mileage would be {}'.format(output))

if __name__ == "__main__":
    app.run()

* Serving Flask app '__main__'
* Debug mode: off

WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead.
* Running on http://127.0.0.1:5000
Press CTRL+C to quit
127.0.0.1 - - [26/Jun/2023 17:50:08] "POST /predict HTTP/1.1" 200 -
127.0.0.1 - - [26/Jun/2023 17:50:08] "GET /static/css/style.css HTTP/1.1" 404 -
```

Predict MVG

<input type="text" value="Cylinders"/>	<input type="text" value="Displacement"/>	<input type="text" value="Horsepower"/>	<input type="text" value="weight"/>	<input type="text" value="Acceleration"/>	<input type="text" value="Model year"/>	<input type="text" value="Origin_2"/>	<input type="text" value="Origin_3"/>	<input type="button" value="Predict"/>
--	---	---	-------------------------------------	---	---	---------------------------------------	---------------------------------------	--

a car mileage would be 17.47